

Spa-Bench: A Real-Robot Benchmark for Spatial Grounding in Vision-Language-Action Models

Anonymous Authors




Fig. 1. **Overview of Spa-Bench.** Left: representative examples from the six spatially grounded tasks: *Object State*, *Relative Size*, *Referential Description*, *Relational Placement*, *Ordinal Position*, and *Counting*. Top right: composition of the 4,092-episode dataset, comprising 1,200 human-teleoperated demonstrations (200 per task family) and 2,892 real-robot evaluation rollouts across five VLA policies. Bottom right: an illustrative matched-scene comparison in which the scene and required physical action remain the same while a direct manipulation instruction is replaced by an OOD spatially grounded instruction. This comparison helps distinguish spatial-grounding failures from an inability to execute the underlying manipulation.

Abstract—Vision-language-action models (VLAs) offer a promising route toward general-purpose robotics by transferring semantic knowledge from large-scale vision-language pretraining into language-conditioned control. However, most VLA evaluations emphasize aggregate manipulation success under changes in objects, instructions, or environments, leaving unclear how reliably these policies ground spatial descriptions in the physical scenes where they act. We therefore ask a sharper question: *how reliably can current VLAs translate spatial descriptions into appropriate real-robot behavior?* To investigate this, we introduce Spa-Bench, a real-robot benchmark that evaluates spatial grounding across six tabletop tasks covering *Object State*, *Relative Size*, *Referential Description*, *Relational Placement*, *Ordinal Position*, and *Counting*. Our primary evaluation compares in-distribution (ID) and out-of-distribution (OOD) spatial configurations while retaining training support for the relevant objects, concepts, and actions, reducing object and motor unfamiliarity as alternative explanations for failure. Four diagnostic conditions further probe physical execution, sensitivity to wording, object-task associations, and recognition that no action is required. We release 4,092 episodes collected on the SO-101 robot arm: 1,200 human-teleoperated demonstrations and 2,892 real-robot evaluation rollouts from five state-of-the-art VLA policies. Across the three fully evaluated policies, success falls from 67.5–73.3% on ID configurations to 40.7–52.0% on OOD configurations, while matched direct instructions improve success by 28.3–40.0 percentage points. These results show that the evaluated VLAs can often execute the required manipulations but lack robust

spatial grounding. Models, datasets, evaluation rollouts, and supporting spreadsheets are available at <https://huggingface.co/datasets/Spa-Bench/spa-bench-data-information> (anonymized).

Index Terms—vision-language-action models, robot learning, spatial grounding, robot evaluation, VLA benchmarking

I. INTRODUCTION

Vision-language-action models (VLAs) have emerged as a promising approach to general-purpose robot learning. By connecting pretrained vision-language representations to robot actions, these models can transfer semantic knowledge acquired from internet-scale data into language-conditioned control [1], [4], [10], [12]. This transfer promises robots that can respond to new objects, instructions, and environments without requiring a separately trained policy for every behavior.

Existing benchmarks have demonstrated increasingly broad manipulation capabilities, including generalization across objects, instructions, visual conditions, and task sequences [5], [6], [20]. Nevertheless, completing a manipulation task does not by itself establish that a policy has correctly interpreted the spatial structure of its environment. Spatial understanding is typically embedded within an overall task-success metric and is consequently entangled with object recognition, language interpretation, and motor execution. If a policy fails to place a green pen to the left of a red block, for example, the failure

could arise from incorrect spatial grounding, failure to identify the objects, an unsuccessful grasp, or an inability to execute the required placement. Even success in a fixed scene may result from recurring object locations or action patterns rather than from grounding the instructed spatial relation.

In this work, we use *spatial grounding* to describe the ability to map language about properties, positions, quantities, and relations in the current physical scene to an appropriate target, goal configuration, or decision not to act. We study this capability through the following question:

How reliably can current VLAs ground spatial descriptions into appropriate actions in the real world?

Answering this question requires more than measuring success on an ordinary held-out task. A test task that simultaneously introduces an unfamiliar object, instruction, spatial concept, or physical behavior cannot reveal which of these changes caused a failure. A diagnostic spatial benchmark should instead establish familiarity with the constituent objects, concepts, and manipulation behaviors before evaluating whether the policy can use the relevant spatial information in a different physical configuration.

We introduce the *Spatial Grounding Benchmark* (Spa-Bench), a controlled real-robot benchmark designed around this principle. Spa-Bench comprises six tabletop tasks evaluating the grounding of *Object State*, *Relative Size*, *Relational Placement*, *Referential Description*, *Ordinal Position*, and *Counting*. All tasks use the same embodiment, workspace, receptacle, and broad pick-and-place action family. The principal source of variation is therefore the scene-grounded rule that determines which object should be manipulated, where it should be placed, or whether movement is even required.

Spa-Bench evaluates policies in ID and OOD spatial configurations. ID configurations measure whether a policy has acquired the spatially grounded behaviors represented in its demonstrations. OOD configurations change the association between a spatial concept and its task argument while preserving training support for the relevant objects, spatial concepts, and physical actions. We implement this controlled split through a *withheld-cell* design: each component of an OOD evaluation configuration appears elsewhere in training, but their particular combination is reserved for evaluation (see Fig. 3). The purpose of this construction is to reduce object and motor unfamiliarity as alternative explanations for failure.

To characterize spatially grounded performance beyond the primary ID–OOD comparison, Spa-Bench includes four additional diagnostic conditions that probe complementary aspects of policy behavior. Direct-manipulation controls remove the need for spatial inference while preserving the requested robot motion. Prompt-paraphrase trials restate both OOD spatial and direct-manipulation instructions on the same scenes, testing whether their performance gap depends on exact wording. Cross-task-object rollouts vary the usual association between an object and task. Finally, no-intervention scenes test whether a policy can recognize when no action is required.

Spa-Bench contains 1,200 human-teleoperated demonstrations, with 200 demonstrations for each task. We evalu-

ate five fine-tuned state-of-the-art policies that are drawn from four different VLA architectures: VLA-0, MolmoAct2, $\pi_{0.5}$, GR00T-N1.7 Frozen LLM (GR00T-N1.7 F-LLM), and GR00T-N1.7 Full Fine-Tune (GR00T-N1.7 F-F). Three policies complete the full 900-rollout evaluation protocol, producing 2,700 evaluation episodes, while two partial baselines contribute 192 episodes. In total, we release 4,092 SO-101 episodes, including the demonstrations, rollout videos, task specifications, and documented scene layouts.

Across the three fully evaluated policies, ID spatial-configuration success ranges from 67.5% to 73.3%, but success decreases to 40.7%–52.0% in OOD spatial configurations. On the OOD spatial configuration matched scene subset, direct-manipulation instructions achieve 81.7%–89.2% success, showing that many of the physical behaviors are executable even when spatially conditioned performance fails. Performance is strongly task dependent: all three policies fail every OOD *Counting* trial, almost always without acting; other tasks are dominated by target-selection or spatial-outcome errors. The direct manipulation advantage also persists when both instruction types are paraphrased, revealing that the state-of-the-art VLAs we evaluate do not exhibit robust spatial grounding. In summary, our core contributions are:

- We introduce Spa-Bench, a six-task real-robot benchmark for evaluating spatial grounding across ID and OOD physical configurations.
- We develop a controlled evaluation protocol that combines a *withheld-cell* split with four diagnostics that probe physical execution, wording sensitivity, object context, and action gating.
- We evaluate five fine-tuned VLA policies and release 4,092 SO-101 episodes, including 1,200 human-teleoperated demonstrations and 2,892 evaluation rollouts, together with task specifications and scene layouts.

II. RELATED WORK

A. Vision-Language-Action Models

Vision-language-action models (VLAs) connect pretrained multimodal representations to robot actions with the goal of transferring semantic and visual knowledge into closed-loop control. RT-2 helped establish this formulation by representing actions as tokens and co-fine-tuning a vision-language model on web and robot data, demonstrating transfer to novel objects, instructions, and semantically specified behaviors [1]. OpenVLA subsequently introduced an open-source 7-billion-parameter VLA trained on large-scale robot data and designed for adaptation to downstream embodiments [2]. In parallel, π_0 coupled a pretrained vision-language backbone to a flow-matching action expert for generating continuous action chunks across multiple embodiments [3].

More recent VLAs have targeted open-world generalization, heterogeneous robot data, embodied reasoning, and practical deployment. $\pi_{0.5}$ extends π_0 through heterogeneous co-training on robot trajectories, verbal instructions, semantic predictions, and web data [4]. The GR00T-N1.7 family uses

a dual-system architecture in which a vision-language module interprets the scene and instruction while a diffusion transformer produces continuous actions [10]. VLA-0 investigates whether continuous actions can instead be generated directly as numerical text without introducing a specialized action vocabulary or action head [11]. MolmoAct2 connects an embodied-reasoning vision-language model to a flow-matching continuous-action expert [12]. Because these models differ in pretraining data, training objectives, action representations, and fine-tuning procedures, Spa-Bench compares complete policy systems rather than isolating individual architectural choices. Each fine-tuned policy is evaluated under the same controlled spatial configurations and diagnostic conditions to characterize its spatial generalization.

B. Simulation and Reasoning Benchmarks

Simulation benchmarks have enabled scalable evaluation of language-conditioned robot behavior. RL-Bench provides a broad collection of vision-guided manipulation tasks and automatically generated demonstrations [14], while ALFRED evaluates grounded instruction following over compositional household activities [15]. LIBERO studies declarative and procedural knowledge transfer across lifelong-learning task suites [16]. RoboCasa expands language-conditioned manipulation to diverse household activities in simulated kitchens [17], and BEHAVIOR-1K targets long-horizon everyday activities involving physical and semantic state changes [18]. VLABench further evaluates semantic knowledge, spatial relationships, physical understanding, and long-horizon reasoning across numerous tasks and objects [5]. VIMA-Bench more explicitly separates forms of generalization, including a combinatorial split that evaluates unseen object–texture pairings whose individual components appear during training [13]. These benchmarks establish task breadth and scalable evaluation, but spatial grounding is generally one component of overall task success and therefore remains entangled with perception, planning, and execution.

A complementary line of work examines whether policy performance persists under controlled distribution shifts. The COLOSSEUM evaluates manipulation policies under systematic visual and physical perturbations [19], while LIBERO-PRO varies manipulated objects, initial states, instructions, and environments, revealing that strong standard-benchmark performance can depend on memorized layouts and action sequences [6]. STAR-Gen organizes generalization along visual, semantic, and behavioral axes [20]. BeTTER applies targeted interventions, including spatial-layout shifts, together with kinematic isolation intended to separate reasoning failures from execution limitations [8]. Other benchmarks evaluate reasoning without requiring closed-loop physical execution. ManipBench probes low-level manipulation reasoning in vision-language models [21], whereas MultihopSpatial evaluates compositional spatial inference and visual grounding over multi-hop relationships [22]. These evaluations provide fine-grained measures of reasoning, but correct visual or textual

answers do not necessarily establish that the inferred relation can guide real-robot behavior.

C. Real-World Robot Evaluation

Real-world benchmarks address the complementary challenges of physical validity, reproducibility, and evaluation at scale. REALM evaluates manipulation policies under controlled perturbations in simulation and validates the resulting trends through physical experiments [7]. RoboChallenge provides infrastructure for evaluating embodied policies across multiple real robots and manipulation tasks [23]. RADAR focuses on real-world dynamics and spatial-physical intelligence while using automated three-dimensional success metrics [9]. ManipArena combines diverse real-robot tasks with task-specific OOD conditions, including unseen objects and held-out spatial, ordering, or procedural configurations of familiar objects [24]. VLA-REPLICA instead emphasizes low-cost cross-laboratory reproducibility through standardized hardware and evaluation on the SO-101 platform [25].

Spa-Bench complements these efforts by isolating spatial grounding under closed-loop real-robot execution. It fixes the robot, workspace, object inventory, and broad manipulation behavior, and places the held-out combination specifically in the mapping between spatial concepts and task arguments. Thus, its OOD configurations preserve training support for the constituent objects, spatial concepts, and physical behaviors. Direct-manipulation controls test whether the relevant manipulation behavior remains executable when the task-specific spatial inference is removed. Prompt paraphrases test whether the performance difference between matched OOD spatial and direct instructions persists under alternative phrasings. Cross-task-object evaluations test sensitivity to familiar object–task associations, while no-intervention scenes test whether policies can correctly withhold action. Together, these diagnostics make Spa-Bench a targeted evaluation of how reliably current VLAs translate spatial descriptions into appropriate behavior in physical scenes.

III. SPA-BENCH DESIGN

A. Task Suite

Spa-Bench evaluates whether a robot policy can ground spatial descriptions in the current physical scene and translate them into an appropriate action. Specifically, the benchmark contains six tabletop tasks spanning three axes of spatial grounding: attribute, relation, and scene structure. The *Object State* and *Relative Size* tasks require selecting an object according to a visually grounded attribute, the *Relational Placement* and *Referential Description* tasks require interpreting directed or object-relative spatial relations whereas the *Ordinal Position* and *Counting* tasks require reasoning about the organization of the scene as a whole.

All six tasks use the same SO-101 robot embodiment, workspace, receptacle, and broad pick-and-place action family. The primary source of variation is therefore the scene-grounded rule that determines which object should be manipulated, where it should be placed, or whether an action



Fig. 2. **Object inventory used in Spa-Bench.** The benchmark uses 27 movable object instances in training and evaluation spanning 11 morphological families. These objects are recombined across the six tasks while the robot embodiment, workspace, and receptacle remain fixed.

is necessary. Five tasks additionally include diagnostic scenes in which no intervention is required because the requested condition already holds or no valid target exists; *Relative Size* is excluded because each comparison set contains a unique smallest and largest object; introducing a size tie would create target ambiguity rather than a clear no-action condition.

All directional and ordinal descriptions are defined from the robot’s perspective. For example, *left*, *right*, *front*, and *behind* refer to directions in the robot’s workspace rather than directions relative to an individual camera view. Table I summarizes the six tasks and their success criteria.

B. Dataset and Physical Setup

The Spa-Bench dataset contains 1,200 human-teleoperated demonstrations, with 200 demonstrations per task, totaling approximately 5.7 hours and 612,733 recorded frames. Each episode contains synchronized robot states, actions, natural-language instructions, and five RGB camera views. The evaluated policies receive the middle and wrist views. Data collection and evaluation are conducted in the same windowless room with fixed lighting and camera placement. We release the demonstrations, prompt-to-episode checklist, scene specifications, initial robot configurations, and rollout videos. The objects used in our dataset are shown in Fig. 2.

C. Controlled ID and OOD Configurations


Existing benchmarks typically evaluate policies on held-out tasks that simultaneously introduce a new object, instruction, spatial concept, or physical behavior. A failure under such a shift cannot be attributed specifically to spatial grounding. Spa-Bench instead constructs OOD spatial configurations that preserve training exposure to the relevant objects, spatial concepts, and manipulation behaviors. During training, the policy encounters every tested spatial concept and sees each evaluation object undergo the required motion. However, selected concept–object or concept–motion pairings never appear with spatially grounded instructions in fine-tuning. OOD evaluation tests these unseen pairings. Because the objects and motions are already familiar from direct instructions or non-revealing spatial descriptions, this design reduces object novelty and motor execution as alternative explanations for failure.

We represent the training support from each task in the fine-tuning dataset as a matrix whose rows denote spatial concepts and whose columns denote task-specific arguments. Concept–argument pairings reserved for OOD evaluation form *withheld*

Object State	Object			
	White Cup	Spice Container	Brown Cup	Striped Cup
Open	Observed	Observed	Observed	Withheld
Closed	Observed	Observed	Withheld	Observed
Upright	Observed	Withheld	Observed	Observed
Upside-down	Withheld	Observed	Observed	Observed



“Move a white cup into the bowl.”



“Move the **upside-down** yellow-white striped cup into the bowl.”

Fig. 3. **Object State concept–argument matrix.** Rows denote object-state concepts and columns denote object identities. *Observed* cells receive state-conditioned instructions during training and are used for ID spatial evaluation, whereas withheld cells receive direct manipulation instructions and are reserved for OOD spatial evaluation. In this example, the *upside-down* white cup is manipulated under a direct instruction, while the concept *upside-down* is paired with the striped cup during training. Hence, the combination of *upside-down* and the white cup is withheld.

cells: the constituent concept, argument, and manipulation behavior appear elsewhere in training, but their particular pairing is absent from the corresponding spatial instructions.

Figure 3 illustrates this construction for the *Object State* task, where rows represent object states and columns represent object identities. Here, the combination of *upside-down* and the white cup is withheld from spatially grounded training instructions. The policy still sees the upside-down white cup manipulated under a direct instruction and receives instructions involving *upside-down* with other objects, including the yellow-white striped cup. OOD evaluation thus pairs a familiar object-state concept with a familiar object while requiring a familiar pick-and-place behavior. Success requires distinguishing the upside-down white cup from an upright white cup and selecting the correct instance, rather than relying on a memorized state–object association.

The same construction applies across tasks, with arguments corresponding to an object, ordered object pair, ordinal slot, or add/remove transition. In *Relational Placement*, for example, training demonstrations move the green pen to the right of the red block under the non-directional instruction “Place the green pen next to the red block,” while *right* is paired with other object pairs. The specific pairing of *right* with the green-pen–red-block pair is reserved for OOD evaluation.

D. Evaluation Protocol and Conditions

A rollout is one autonomous policy execution from a specified scene. Each fully evaluated policy completes 900 physical rollouts across the six tasks. Per task, the primary evaluation contains 20 ID and 50 OOD spatial-configuration rollouts.

TABLE I
SPA-BENCH TASKS. ALL TASKS USE THE SAME ROBOT, WORKSPACE, RECEPTACLE, AND BROAD PICK-AND-PLACE BEHAVIOR WHILE VARYING THE SPATIAL INFORMATION NEEDED TO DETERMINE THE CORRECT RESPONSE.

Task	Grounding Axis	Example Instruction	Success Criterion
Object State	Attribute	“Place the <i>closed</i> white cup in the bowl.”	The same-category object whose object state matches the instruction is placed in the bowl.
Relative Size	Attribute	“Place the <i>smallest</i> object in the bowl.”	The object satisfying the requested relative-size description is placed in the bowl.
Relational Placement	Relation	“Place the green pen to the <i>left of</i> the red block.”	The target finishes within 3 inches of the reference and within the requested $\pm 45^\circ$ directional sector.
Referential Description	Relation	“Place the pencil sharpener <i>furthest</i> from the yellow pen in the bowl.”	The object satisfying the specified relation to the reference object is placed in the bowl.
Ordinal Position	Scene structure	“Place the <i>third</i> object <i>from the left</i> in the bowl.”	The object occupying the requested ordinal position is placed in the bowl.
Counting	Scene structure	“Leave exactly <i>two</i> objects in the bowl.”	The requested number of objects remains in the bowl at the end of the episode.

ID configurations evaluate concept–argument pairings present in spatially grounded training instructions (*observed cells*), whereas OOD configurations evaluate pairings absent from those instructions despite training exposure to the constituent concepts, objects, and motions (*withheld cells*).

For 20 of the 50 OOD scene specifications per task, we repeat the evaluation with a direct-manipulation instruction that removes the need for task-specific spatial inference while preserving the scene setup and target object. These control instructions also occur for all objects during fine-tuning, but less frequently than the spatially grounded instructions. For example, the *Referential Description* instruction “Put the crushed soda can closest to the airplane toy with propellers in the bowl” requires resolving a relational referent. Its direct control, “Move the crushed Coke can into the bowl,” requests the same robot motion but names the target explicitly (Fig. 4). A direct-control success paired with a spatial wrong-instance failure therefore provides evidence that referent resolution, rather than pick-and-place execution alone, caused the failure.

For 10 of these 20 matched scenes per task, we also provide paraphrased versions of both instruction types, producing four conditions: original OOD spatial, paraphrased OOD spatial, original direct manipulation, and paraphrased direct manipulation. Within each original–paraphrase pair, the scene setup, intended target, and instruction meaning remain the same while the phrasing changes. This measures wording sensitivity and tests whether the spatial–direct performance difference persists under alternative phrasings (Also shown in Fig. 4).

An additional 20 rollouts per task evaluate whether ID spatial grounding transfers to objects normally associated with another task. These cross-task-object rollouts retain a familiar spatial concept and manipulation behavior while changing the object context. Because no object from another task instantiates all required *Object State* conditions, that task instead uses a separate novel-object pilot and is excluded from the strict cross-task-object aggregate.

Five tasks additionally include 20 no-intervention rollouts. In *Counting* and *Relational Placement*, the requested count or relation already holds. In *Object State*, *Referential Description*, and *Ordinal Position*, the instruction has no valid target in the




Fig. 4. **Matched-scene diagnostics.** The same scene setup and intended object motion are evaluated under four instruction conditions: OOD spatial, direct manipulation, and a paraphrase of each. In this *Referential Description* example, the direct instruction names the target explicitly, whereas the spatial instruction requires resolving it from the scene.

scene. Success requires initiating no arm or gripper motion. *Relative Size* omits this condition because every comparison set contains a unique smallest and largest object. *Relational Placement* also includes 20 matched movement-required controls to distinguish recognition that a relation already holds from a general tendency to remain stationary.

These task-specific adjustments produce 170 rollouts for *Relational Placement*, 130 for *Relative Size*, and 150 for each remaining task, totaling 900 rollouts per fully evaluated policy.

E. Cross-Policy Scene Matching, Scoring, and Analysis

To support comparisons across policies, corresponding rollouts are manually reset to the same scene specification (e.g., evaluation 103 for MolmoAct2 uses the layout specified for evaluation 103 for $\pi_{0.5}$). Resets reproduce object identities and locations using reference scene images; pixel coordinates

are not matched. Within each task, the 20 direct-manipulation rollouts reuse a subset of the 50 OOD scene specifications, and 10 of the direct-manipulation scene specifications are additionally reused for both paraphrase conditions.

The primary outcome is binary physical task success defined in Table I; partial credit is not assigned. Rollouts are limited to 60 seconds for all models except VLA-0, which receives 90 seconds due to its slower inference rate. No human intervention occurs after the rollout timer and scoring begin. We report success rates with two-sided Wilson 95% confidence intervals [27]. Matched-scene differences are reported as descriptive paired contrasts computed only over the corresponding reused scene subsets. Qualitative failure annotations are assigned separately and do not alter binary scoring.

IV. MODELS AND EXPERIMENTAL SETUP

We evaluate five VLAs from four representative architectures: VLA-0 [11], MolmoAct2 [12], $\pi_{0.5}$ [4], and two GR00T-N1.7 adaptations [10]. These architectures span textual [11] and continuous [4], [10], [12] action representations, with action-chunk horizons from 8 to 50 steps [4], [10]–[12]. Because these factors vary across architectures, cross-architecture performance differences cannot be attributed to any single design choice. The two GR00T-N1.7 adaptations provide a narrower within-architecture comparison of fine-tuning strategies. GR00T-N1.7 F-LLM freezes only the language-model parameters; it continues to receive language instructions while its visual, multimodal, and action-generation modules are updated. GR00T-N1.7 F-FT updates all parameters.

A. Training and Deployment Summary

Each policy receives the six-dimensional robot state alongside middle and wrist-camera observations. All policies are trained for 12 epochs with a global batch size of 96 on the same 1,200 episode identities using their native preprocessing and fine-tuning procedures. To suppress high-frequency action jitter, both GR00T-N1.7 variants use a causal action-smoothing filter during deployment [26].

B. Evaluation Scope

MolmoAct2, $\pi_{0.5}$, and GR00T-N1.7 F-LLM complete all 900 rollouts and support the full ID–OOD and diagnostic analyses. GR00T-N1.7 F-FT completes only the 120-rollout ID condition, while VLA-0 completes 72 ID rollouts; both are reported as partial ID baselines. Across the five policies, the evaluation contains 2,892 real-robot rollouts. Diagnostic conditions are analyzed separately, with the matched-control gap computed as direct-manipulation success minus OOD success on the corresponding 120 matching scenes. Prompt-paraphrase results are separated into paraphrases of OOD spatial instructions and paraphrases of direct manipulation instructions. The strict cross-task-object aggregate excludes the separate 20-trial *Object State* task novel-object pilot. No-intervention performance is reported by task because the semantic reason for not acting differs across tasks.

V. EXPERIMENTAL RESULTS

A. Spatial Grounding: ID and OOD Configurations

ID spatial grounding first establishes whether each policy acquired the training-supported spatial behaviors. MolmoAct2 achieved 88/120 successes (73.3%, Wilson 95% CI: 64.8–80.4%), $\pi_{0.5}$ achieved 82/120 (68.3%, 59.6–76.0%), and GR00T-N1.7 F-LLM achieved 81/120 (67.5%, 58.7–75.2%). Because neither partial baseline established adequate ID task competence (25/120 for GR00T-N1.7 F-FT and 0/72 for VLA-0), their OOD failures would not distinguish failed recombination from failure to acquire the underlying tasks; both are therefore retained only as diagnostic ID baselines. Fig. 5 reports the task-level ID results for all five policies.

Performance decreased on the configurations withheld by the *withheld-cell* split. Across 300 OOD spatial-grounding rollouts per policy, MolmoAct2 achieved 156/300 successes (52.0%, Wilson 95% CI: 46.4–57.6%), GR00T-N1.7 F-LLM achieved 133/300 (44.3%, 38.8–50.0%), and $\pi_{0.5}$ achieved 122/300 (40.7%, 35.3–46.3%). Performance varied by task rather than producing a consistent model ranking: GR00T-N1.7 F-LLM performed best on *Relational Placement*, whereas MolmoAct2 performed best on *Object State* and *Referential Description*. All three policies failed every OOD *Counting* rollout (Fig. 6), showing that familiarity with objects and physical transitions did not make the withheld count-conditioned combinations usable. Relative to ID, this decrease indicates incomplete recombination of familiar spatial concepts and actions rather than a uniform loss of task competence.

B. Matched-Scene Diagnostics

To determine whether the OOD decrease reflects spatial understanding rather than robot execution, 20 OOD scene specifications per task are repeated under direct-manipulation instructions, yielding 120 matched scenes per policy. MolmoAct2 achieved 73/120 successes (60.8%) under OOD spatial instructions and 107/120 (89.2%) under direct manipulation, a 28.3-percentage-point difference. The corresponding results were 54/120 (45.0%) vs. 100/120 (83.3%) for $\pi_{0.5}$, a 38.3-point difference, and 50/120 (41.7%) vs. 98/120 (81.7%) for GR00T-N1.7 F-LLM, a 40.0-point difference. Because these comparisons preserve the scene setup and relevant manipulation behavior while removing the task-specific spatial concept, the consistently higher direct-control success shows that execution difficulty alone does not account for the OOD spatial-grounding failures (Fig. 6).

To test whether this contrast depends on exact wording, 10 matched scenes per task also receive paraphrased versions of both instruction types. On this 60-scene subset, spatial success decreased from 48.3% to 38.3% for GR00T-N1.7 F-LLM, from 48.3% to 41.7% for $\pi_{0.5}$, and from 65.0% to 50.0% for MolmoAct2. Under direct-manipulation instructions, success decreased from 81.7% to 70.0% for GR00T-N1.7 F-LLM and remained at 86.7% for $\pi_{0.5}$ and 90.0% for MolmoAct2. Thus, paraphrasing consistently reduced spatial success, whereas its effect on direct instructions was model dependent (Fig. 7).




Fig. 5. **Task-level ID spatial-grounding success.** Bars show per-task success rates, and whiskers indicate two-sided Wilson 95% confidence intervals. Complete ID evaluations contain 20 trials per task (120 total); VLA-0 contains 72 trials, with at least 10 per task. GR00T-N1.7 F-FT and VLA-0 are reported only as diagnostic ID baselines because their low ID success would make subsequent OOD failures difficult to interpret.




Fig. 6. **Task-level OOD spatial-grounding and direct-manipulation success.** Colors denote policies; dark and light bars denote OOD spatial and direct-manipulation instructions, respectively. Whiskers show two-sided Wilson 95% confidence intervals. OOD bars use 50 trials per task and direct bars use 20; the matched comparison in Sec. V-B uses only the corresponding 20 OOD scenes.




Fig. 7. **Prompt-paraphrase diagnostic on 60 matched scenes.** Dark bars show original instructions; light bars show paraphrased instructions for the same scenes.

C. Additional Diagnostics

Cross-task-object trials test whether ID spatial grounding transfers to objects from other tasks. Across five comparable tasks, success fell from 64/100 to 55/100 for GR00T-N1.7 F-LLM, from 66/100 to 54/100 for $\pi_{0.5}$, and from 70/100 to 53/100 for MolmoAct2. Because the conditions cover the same tasks but are not paired scenes differing only in object identity, these decreases are descriptive. Their consistent direction nevertheless suggests sensitivity to object-task associations learned during fine-tuning. In the *Relational Placement* trials, none of the three policies defaulted to placing the juggling ball in the bowl, although every fine-tuning demonstration that manipulated the ball placed it in the bowl. This suggests some transfer beyond the learned object-goal association, although the requested relation was not always satisfied.

Action gating was also strongly task dependent. Across the four no-intervention diagnostics without matched corrective controls, aggregate success was only 20.0–23.8%. Success reached 80–85% when the *Counting* goal was already satisfied and 0% for both *Object State* and *Ordinal Position* across all three policies tested. This contrast shows that the policies can remain stationary but do not reliably condition that decision

on whether the instruction already holds or has a valid referent. We analyze *Relational Placement* separately because its no-movement scenes are paired with corresponding scenes requiring corrective movement. Success was 75–95% when the requested relation already held but only 5–15% when correction was required. The reversed performance gap suggests that the high no-movement success primarily reflects a task-specific tendency to remain stationary when the target and reference objects are nearby, rather than reliable recognition of whether intervention is required.

D. Analysis of Failures by Task

For each failed rollout, we recorded its behavior and assigned a mutually exclusive primary failure mode from the taxonomy in Fig. 8. When multiple problems occurred, we assigned the mode judged most directly responsible for failure. OOD patterns are strongly task dependent: no action accounts for 148/150 *Counting* failures, wrong object selection for 74/79 *Ordinal Position* failures, and same-category, wrong-instance selection for 75/77 *Object State* failures. In contrast, most *Relational Placement* failures from $\pi_{0.5}$ and MolmoAct2 involve incorrect spatial outcomes after approaching the intended object (32/39), while GR00T-N1.7 F-LLM exhibits a more heterogeneous mixture. These results highlight different bottlenecks across tasks and policies rather than a single failure mode that explains performance across the benchmark.

VI. CONCLUSION AND LIMITATIONS

We introduced Spa-Bench to evaluate spatial grounding across six real-robot tasks. Across the three fully evaluated VLAs, success decreased from ID to OOD spatial configurations. On matched scenes, performance recovered when the target was named directly, and this advantage persisted under paraphrasing. Cross-task objects revealed sensitivity to learned




Fig. 8. **Primary Failure-Mode Composition for OOD Spatial Configurations.** Bars are normalized within each task–model pair, and n denotes the number of failures. Wrong-instance errors are separated from other wrong-target errors only for *Object State* and *Referential Description*.

object associations. Failure modes were task-dependent, with some patterns consistent across VLAs. These findings show that successful manipulation does not necessarily imply spatial understanding. Spa-Bench’s *withheld-cell* design provides a controlled testbed for measuring whether VLAs select the correct behavior based on the spatial information expressed in an instruction: a capability essential for reliable autonomy. **Limitations.** Spa-Bench deliberately fixes the robot embodiment, workspace, object inventory, receptacle, and manipulation family. This enables targeted attribution but limits external validity: scenes contain at most five candidate objects, resets are not pixel-identical, and the no-intervention conditions are zero-shot. Model-specific preprocessing and adaptation choices also mean that the results characterize complete policy systems rather than individual architectural components.

REFERENCES

- [1] B. Zitkovich et al., “RT-2: Vision-language-action models transfer web knowledge to robotic control,” in *Proc. CoRL*, pp. 2165–2183, 2023.
- [2] M. J. Kim et al., “OpenVLA: An open-source vision-language-action model,” in *Proc. CoRL*, pp. 2679–2713, 2025.
- [3] K. Black et al., “ π_0 : A vision-language-action flow model for general robot control,” in *Proc. RSS*, 2025.
- [4] K. Black et al., “ $\pi_{0.5}$: A vision-language-action model with open-world generalization,” in *Proc. CoRL*, pp. 17–40, 2025.
- [5] S. Zhang et al., “VLABench: A large-scale benchmark for language-conditioned robotics manipulation with long-horizon reasoning tasks,” in *Proc. ICCV*, pp. 11142–11152, 2025.
- [6] X. Zhou et al., “LIBERO-PRO: Towards robust and fair evaluation of vision-language-action models beyond memorization,” arXiv:2510.03827, 2025.
- [7] M. Sedlacek et al., “REALM: A real-to-sim validated benchmark for generalization in robotic manipulation,” *IEEE Robot. Autom. Lett.*, vol. 11, no. 7, pp. 8315–8322, 2026.
- [8] H. Xu, S. Zheng, H. Luo, W. Zhang, Z. Xi, and Z. Lu, “Unmasking the illusion of embodied reasoning in vision-language-action models,” in *Proc. ECCV*, 2026, to appear.
- [9] Y. Chen et al., “RADAR: Benchmarking vision-language-action generalization via real-world dynamics, spatial-physical intelligence, and autonomous evaluation,” arXiv:2602.10980, 2026.
- [10] NVIDIA et al., “GR00T N1: An open foundation model for generalist humanoid robots,” arXiv:2503.14734, 2025.
- [11] A. Goyal, H. Hadfield, X. Yang, V. Blukis, and F. Ramos, “VLA-0: Building state-of-the-art VLAs with zero modification,” arXiv:2510.13054, 2025.
- [12] H. Fang et al., “MolmoAct2: Action reasoning models for real-world deployment,” in *Proc. CoRL*, 2026, to appear.
- [13] Y. Jiang et al., “VIMA: Robot manipulation with multimodal prompts,” in *Proc. ICML*, pp. 14975–15022, 2023.
- [14] S. James, Z. Ma, D. Rovick Arrojo, and A. J. Davison, “RLBench: The robot learning benchmark & learning environment,” *IEEE Robot. Autom. Lett.*, vol. 5, no. 2, pp. 3019–3026, 2020.
- [15] M. Shridhar et al., “ALFRED: A benchmark for interpreting grounded instructions for everyday tasks,” in *Proc. CVPR*, pp. 10740–10749, 2020.
- [16] B. Liu et al., “LIBERO: Benchmarking knowledge transfer for lifelong robot learning,” in *Proc. NeurIPS*, vol. 36, pp. 44776–44791, 2023.
- [17] S. Nasiriany et al., “RoboCasa: Large-scale simulation of everyday tasks for generalist robots,” in *Proc. RSS*, 2024.
- [18] C. Li et al., “BEHAVIOR-1K: A benchmark for embodied AI with 1,000 everyday activities and realistic simulation,” in *Proc. CoRL*, pp. 80–93, 2023.
- [19] W. Pumacay, I. Singh, J. Duan, R. Krishna, J. Thomason, and D. Fox, “THE COLOSSEUM: A benchmark for evaluating generalization for robotic manipulation,” in *Proc. RSS*, 2024.
- [20] J. Gao, S. Belkale, S. Dasari, A. Balakrishna, D. Shah, and D. Sadigh, “A taxonomy for evaluating generalist robot manipulation policies,” *IEEE Robot. Autom. Lett.*, vol. 11, no. 3, pp. 3182–3189, 2026.
- [21] E. Zhao et al., “ManipBench: Benchmarking vision-language models for low-level robot manipulation,” in *Proc. CoRL*, pp. 3413–3462, 2025.
- [22] Y. Lee, S. Jang, Y. Cho, S. Lee, Y.-J. Lee, and S. J. Hwang, “Multi-hopSpatial: Multi-hop compositional spatial reasoning benchmark for vision-language models,” in *Proc. ECCV*, 2026, to appear.
- [23] A. Yakefu et al., “RoboChallenge: Large-scale real-robot evaluation of embodied policies,” arXiv:2510.17950, 2025.
- [24] Y. Sun et al., “ManipArena: Comprehensive real-world evaluation of reasoning-oriented generalist robot manipulation,” arXiv:2603.28545, 2026.
- [25] A. S. Huang, J. Zhang, S. Tang, and Y. Xiang, “VLA-REPLICA: A low-cost, reproducible benchmark for real-world evaluation of vision-language-action models,” arXiv:2605.20774, 2026.
- [26] ANR Robot, “Experiments with GR00T,” [Online]. Available: <https://www.anrrobot.com/robots/experiments/gr00t.html>. Accessed: Sep. 13, 2026.
- [27] E. B. Wilson, “Probable inference, the law of succession, and statistical inference,” *J. Amer. Statist. Assoc.*, vol. 22, no. 158, pp. 209–212, 1927.